Chun-Yen Pan

AI Systems | LLM Agents | Distributed Systems | Infrastructure
77geo5rge6@gmail.com | yenslife.top | linkedin.com/in/yenslife | github.com/yenslife

Education

National Cheng Kung University **Sep. 2025 – Expected Aug. 2027**
M.S. in Computer Science and Information Engineering *Tainan, Taiwan*
- ISMP Lab; Advisor: Distinguished Prof. Yau-Hwang Kuo. Research interests: LLM Security, Agent Security, and AI Systems.

National Cheng Kung University **Sep. 2021 – Jun. 2025**
B.S. in Computer Science and Information Engineering; GPA: 3.92 / 4.3 *Tainan, Taiwan*

Experience

National Cheng Kung University – ISMP Lab **Feb. 2025 – Present**
Research Assistant *Tainan, Taiwan*
- Contribute to a multi-site distributed infrastructure platform for a major telecommunications operator, including a 6-host environment with 30+ KVM VMs across Linux, networking, Kafka, MongoDB, and high-availability systems; took ownership of VM provisioning, service deployment, automation, and troubleshooting.
- Diagnosed a Kafka consumer throughput bottleneck in a deployment with 120 partitions and 40 consumer instances; redesigned the processing path with asynchronous execution so consumers could keep pace with production workload.
- Investigated partial failover across a 7-VIP Keepalived/VRRP service group and designed a coordinator-based health architecture to enforce all-or-nothing group failover consistency.
- Lead a 3-person student team in an industry–academia AI cybersecurity R&D project developing an LLM-powered adaptive cyber deception system, coordinating technical design, implementation, evaluation, and review deliverables.
- Collaborate with a nursing supervisor at National Cheng Kung University Hospital to build a Dify-based LINE chatbot for dialysis patient education; translated domain requirements and healthcare materials into a controlled workflow serving approximately 50 registered users.

AI Next Mine **Jul. 2024 – Jun. 2025**
AI Engineering Intern *Remote, Taiwan*
- Developed LLM workflows and RAG pipelines for an AI-powered ESG and sustainability reporting assistant using Dify, Langflow, and Python.
- Maintained client-facing AI applications across FastAPI/Flask backends and an existing Next.js frontend, resolving 10+ reported bugs and delivering 15 features within a 3-month period.
- Deployed AI applications on AWS with Docker Compose and selected EC2 configurations based on model resource requirements and operating cost.

EMILY.RPA – Automate Your Daily Work **Jul. 2025 – Aug. 2025**
RPA Coach Intern *New Taipei City, Taiwan*
- Assisted enterprise users in translating repetitive business processes into end-to-end browser automation and RPA workflows, including implementation support and troubleshooting.

Publication & Award

- **Author**, 《零基礎玩轉 LLM 應用全攻略》, 博碩文化, Feb. 2026 – a hands-on guide to building practical LLM applications with Python.
- **1st Place, Generative AI Category**, iThome Ironman 2024 – 30-day technical series on building AI/LLM applications with Python.

Technical Skills

Languages: Python, C/C++, Bash, JavaScript/TypeScript, SQL

AI / LLM: LLM Agents, RAG, Dify, LangChain, Langflow, LoRA fine-tuning, LLM/Agent Security, OpenAI API

Systems / Data: Linux, KVM/QEMU, libvirt, systemd, Docker, Keepalived/VRRP, SR-IOV, Apache Kafka, MongoDB, PostgreSQL

Backend / Cloud: FastAPI, Flask, REST APIs, AWS, Git, GitHub Actions, Docker Compose

AI-assisted Engineering: Codex and ChatGPT for implementation, debugging, technical research, and documentation with manual verification